Abstract

Heart disease is one of the leading causes of mortality worldwide, and early risk identification plays a crucial role in preventing severe health complications. Traditional diagnostic methods often rely on extensive medical tests and expert evaluation, which may not always be accessible or time-efficient.

This project presents a machine learning-based heart disease risk prediction system using patient health parameters. Various classification models, including Logistic Regression, Decision Tree, and Random Forest, were trained and evaluated. Logistic Regression demonstrated the best overall performance and was selected as the final model. The system predicts the probability of heart disease and categorizes patients into Low, Medium, or High risk levels for better interpretability. A Streamlit-based interface was developed to provide real-time prediction, making the system user-friendly and practical for decision support.

Keywords: Heart Disease Prediction, Machine Learning, Logistic Regression, Risk Assessment, Healthcare Analytics

CHAPTER 1

Introduction

Heart disease remains a major global health concern, accounting for a significant number of deaths each year. Factors such as high blood pressure, cholesterol levels, lifestyle habits, and genetic predisposition contribute to the increasing prevalence of cardiovascular diseases. Early detection and preventive measures are essential to reduce mortality rates and improve patient outcomes.

With the advancement of data-driven technologies, machine learning has emerged as an effective tool in healthcare for predicting diseases based on historical patient data. Machine learning models can identify complex patterns and relationships among multiple health parameters that may not be easily detected through traditional diagnostic methods.

This project focuses on developing a heart disease risk prediction system using machine learning techniques. The system analyzes patient health attributes and predicts the likelihood of heart disease. Additionally, it categorizes the predicted probability into different risk levels to enhance interpretability. The objective of this work is to assist in early risk assessment and provide a simple decision-support tool rather than a definitive medical diagnosis.

The report is organized as follows: Chapter 2 discusses the basic concepts and related work, Chapter 3 explains the problem statement and system design, Chapter 4 describes the implementation and results, and Chapter 5 concludes the report with future scope.

CHAPTER 2

Basic Concepts and Literature Review

This chapter discusses the fundamental concepts related to machine learning techniques used in healthcare applications, particularly for heart disease prediction. It also presents a brief review of existing research work in this domain.

2.1 Machine Learning in Healthcare

Machine learning has gained significant attention in the healthcare sector due to its ability to analyze large volumes of medical data and assist in clinical decision-making. By learning patterns from historical patient data, machine learning models can help predict diseases, identify risk factors, and support early diagnosis.

In cardiovascular healthcare, machine learning techniques are widely used for predicting heart disease based on patient attributes such as age, blood pressure, cholesterol levels, and lifestyle-related factors. These predictive systems aim to assist medical professionals by providing risk assessments rather than replacing clinical judgment.

2.2 Heart Disease Prediction Using Machine Learning

Heart disease prediction is typically formulated as a binary classification problem, where the objective is to determine whether a patient is at risk of developing heart disease. Various machine learning algorithms have been applied to this problem, including Logistic Regression, Decision Trees, Support Vector Machines, and ensemble methods.

Most existing systems rely on structured datasets containing patient medical attributes. Feature preprocessing, such as normalization and encoding, plays a crucial role in improving prediction performance. Evaluation metrics like accuracy, precision, recall, and F1-score are commonly used to assess model effectiveness.

2.3 Logistic Regression

Logistic Regression is a widely used supervised learning algorithm for binary classification problems. It estimates the probability of an outcome using a logistic (sigmoid) function and assigns class labels based on a threshold value.

In healthcare applications, Logistic Regression is preferred due to its simplicity, interpretability, and robustness. The model provides probability scores that can be easily interpreted as risk levels, making it suitable for medical decision-support systems where transparency is essential.

2.4 Decision Tree and Random Forest

Decision Trees are classification models that represent decisions in a tree-like structure. They are easy to interpret and can handle both numerical and categorical data. However, Decision Trees are prone to overfitting, especially when trained on small datasets.

Random Forest is an ensemble learning technique that combines multiple Decision Trees to improve prediction accuracy and reduce overfitting. While Random Forest models often achieve higher accuracy, they may sacrifice interpretability, which is an important consideration in healthcare-related applications.

2.5 Data Preprocessing and Feature Scaling

Data preprocessing is a critical step in machine learning pipelines. It includes handling missing values, encoding categorical variables, and scaling numerical features. Feature scaling ensures that all input attributes contribute equally to model training and prevents dominance by features with larger numerical ranges.

Standardization techniques such as StandardScaler are commonly used to normalize data before training machine learning models. Proper preprocessing improves model stability and generalization performance.

CHAPTER 3

Problem Statement and System Design

This chapter describes the problem addressed by the project, the system requirements, and the overall design of the proposed solution.

3.1 Problem Statement

Heart disease is one of the leading causes of death worldwide. Early identification of individuals at risk is essential for timely medical intervention and preventive care. However, traditional diagnostic approaches often require extensive medical tests and expert evaluation, which may not always be accessible or efficient.

The objective of this project is to develop a machine learning-based system that predicts the risk of heart disease using patient health parameters. The system aims to classify individuals into different risk categories and provide an interpretable output that can assist in early risk assessment. The proposed solution is intended to act as a decision-support tool rather than a replacement for medical professionals.

—

3.2 Requirement Specifications

3.2.1 Functional Requirements

The functional requirements of the proposed system are as follows:

- The system should accept patient health parameters as input.
- The system should preprocess the input data before prediction.

- The system should predict the probability of heart disease using a trained machine learning model.
- The system should classify the prediction into Low, Medium, or High risk levels.
- The system should display the prediction result in a user-friendly interface.

3.2.2 Non-Functional Requirements

The non-functional requirements of the system include:

- The system should provide predictions with reasonable accuracy.
- The system should be easy to use and understand.
- The system should produce results within a short response time.
- The system should maintain data consistency during preprocessing and prediction.

—

3.3 System Design

3.3.1 Design Constraints

The system was developed under the following constraints:

- Programming Language: Python
- Development Environment: Anaconda, Jupyter Notebook
- Libraries Used: NumPy, Pandas, scikit-learn, Streamlit
- Platform: Desktop/Laptop environment

3.3.2 System Architecture

The overall architecture of the proposed system follows a modular approach. The system consists of data preprocessing, machine learning model training, evaluation, and user interaction components.

The flow of the system is as follows:

1. User provides health-related input parameters through the interface.
2. Input data is preprocessed and scaled using a trained scaler.
3. The processed data is passed to the trained Logistic Regression model.
4. The model predicts the probability of heart disease.
5. The probability is converted into a risk category (Low, Medium, or High).
6. The final prediction and risk level are displayed to the user.

This structured architecture ensures clarity, modularity, and ease of maintenance.

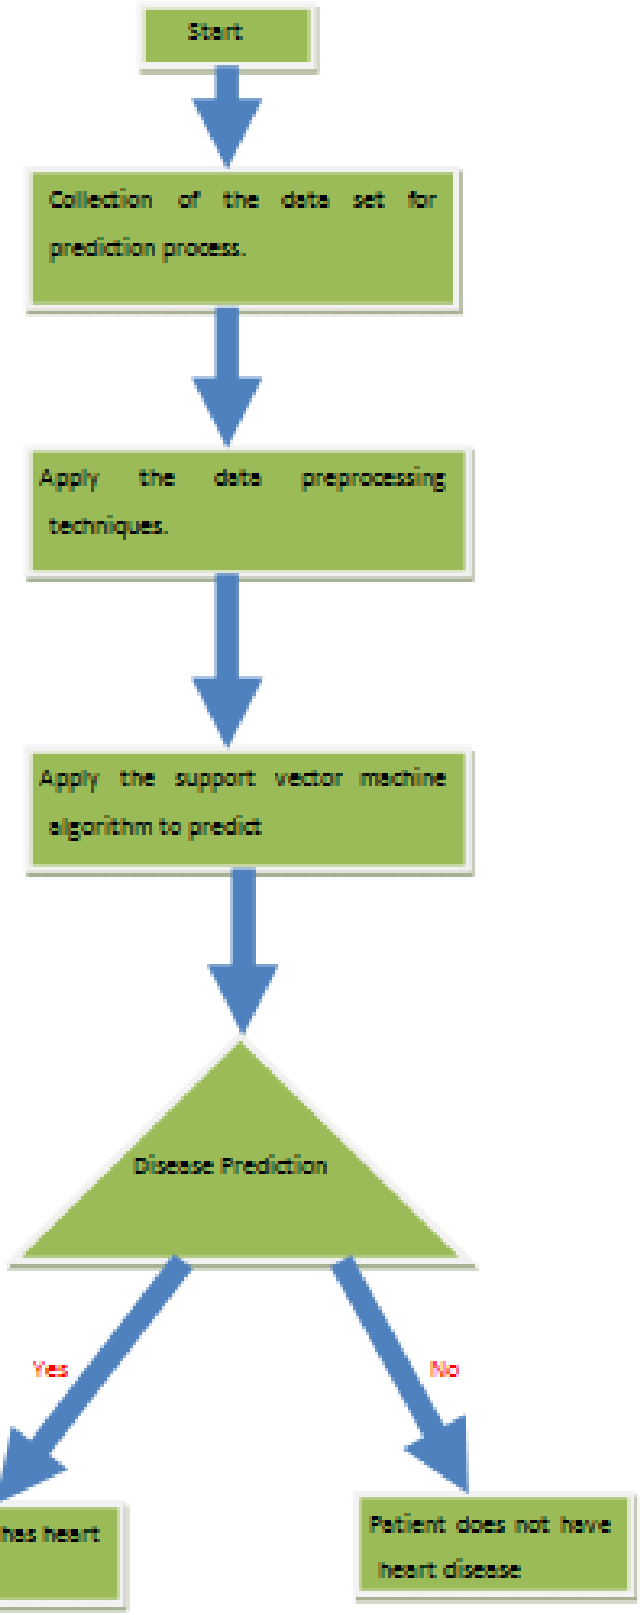

Figure 3.1: System Architecture of Heart Disease Risk Prediction System

CHAPTER 4

Implementation and Result Analysis

This chapter describes the implementation details of the proposed heart disease risk prediction system, including data preprocessing, model training, evaluation, and result analysis.

4.1 Dataset Description

The dataset used in this project was obtained from the UCI Machine Learning Repository. It consists of 270 patient records with 14 attributes. The dataset includes medical parameters such as age, blood pressure, cholesterol level, heart rate, and other clinical indicators relevant to heart disease prediction.

The target variable represents the presence or absence of heart disease. All attributes were found to be complete with no missing values, making the dataset suitable for machine learning-based analysis.

—

4.2 Data Preprocessing

Data preprocessing is a crucial step in the machine learning pipeline. The following preprocessing steps were performed:

- The target variable was encoded into binary form, where 0 represents absence of heart disease and 1 represents presence of heart disease.
- Input features were separated from the target variable.
- Feature scaling was applied using the StandardScaler technique to normalize the data and ensure uniform contribution of all attributes.

- The dataset was divided into training and testing sets using an 80:20 split ratio.

These steps ensured improved model performance and reduced bias caused by varying feature scales.

—

4.3 Model Implementation

Three machine learning classification models were implemented and evaluated in this project:

4.3.1 Logistic Regression

Logistic Regression was used as a baseline classification model due to its simplicity and interpretability. The model estimates the probability of heart disease using a sigmoid function and classifies the output based on a predefined threshold. Logistic Regression is particularly suitable for medical applications due to its transparent decision-making process.

4.3.2 Decision Tree Classifier

The Decision Tree classifier was implemented to model complex decision boundaries. It represents classification decisions in a tree-like structure, making it easy to visualize. However, Decision Trees are prone to overfitting, especially when trained on smaller datasets.

4.3.3 Random Forest Classifier

Random Forest is an ensemble learning technique that combines multiple Decision Trees to improve prediction accuracy and robustness. While Random Forest models often perform well in classification tasks, they may reduce interpretability compared to simpler models.

—

4.4 Model Evaluation Metrics

The performance of the models was evaluated using standard classification metrics:

- Accuracy: Measures the overall correctness of predictions.

- Precision: Indicates the proportion of correctly predicted positive cases.
- Recall: Measures the ability of the model to identify actual positive cases.
- F1-score: Harmonic mean of precision and recall.

These metrics provide a comprehensive evaluation of model effectiveness, particularly in medical risk prediction scenarios where recall is a critical factor.

—

4.5 Result Analysis

The evaluation results of the three models are summarized in Table 4.1.

Table 4.1: Model Performance Comparison

Model	Accuracy	Precision	Recall	F1-Score
Logistic Regression	0.907	0.900	0.857	0.878
Decision Tree	0.685	0.577	0.714	0.638
Random Forest	0.796	0.778	0.667	0.718

Among the three models, Logistic Regression achieved the highest accuracy and F1-score, along with strong recall performance. Since recall is a critical metric in medical risk prediction to minimize false negatives, Logistic Regression was selected as the final model.

—

4.6 Risk Level Classification

To enhance interpretability, the probability output of the Logistic Regression model was converted into three risk categories:

- Low Risk: Probability less than 0.3
- Medium Risk: Probability between 0.3 and 0.6
- High Risk: Probability greater than 0.6

This risk-based classification enables easier understanding of predictions and supports early decision-making.

—

4.7 System Interface

A Streamlit-based graphical user interface was developed to demonstrate real-time heart disease risk prediction. The interface allows users to input patient health parameters and displays the predicted risk level along with the probability score. Screenshots of the interface are included to illustrate system functionality.

CHAPTER 5

Conclusion and Future Scope

5.1 Conclusion

This project successfully developed a machine learning-based heart disease risk prediction system using patient health parameters. Various classification models, including Logistic Regression, Decision Tree, and Random Forest, were implemented and evaluated to determine the most suitable approach.

Based on performance metrics such as accuracy, precision, recall, and F1-score, Logistic Regression demonstrated the best overall performance and was selected as the final model. The system predicts the probability of heart disease and categorizes patients into Low, Medium, and High risk levels, making the output more interpretable and useful for decision support.

The implementation of a Streamlit-based interface further enhances usability by allowing real-time risk prediction through an intuitive user interface. The proposed system can assist in early risk assessment and support preventive healthcare measures.

5.2 Future Scope

The current system can be further enhanced in several ways:

- Incorporating larger and more diverse datasets to improve model generalization.
- Exploring advanced machine learning and deep learning techniques for improved prediction accuracy.
- Integrating real-time clinical data from hospitals and healthcare centers.
- Developing a mobile or web-based application for wider accessibility.

- Adding explainable AI techniques to improve transparency and trust in predictions.

These improvements can help transform the system into a more robust and scalable healthcare decision-support tool.